

## Week 4: Distributed and Cloud Computing

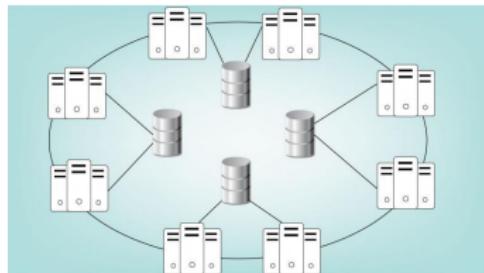
Tien-Nam Le

`tien-nam.le@ens-lyon.fr`  
`perso.ens-lyon.fr/tien-nam.le/su`

Sciences U Lyon

# Outline

## 1. Distributed Computing



## 2. Cloud Computing



# Part I. Distributed Computing

Two distributed computing frameworks



VS



# Problems

## **Problem 1 – SUM() SQL query:**

An international fashion company has many shops, each shop keep data of every sale. Objective: query total revenue from sale of each product.

**Naive solution:** Run through the whole database one-by-one record.

## **Problem 2 – Machine learning problems**

Eg: Linear regression,  $k$ -nearest-neighbor of a large dataset.

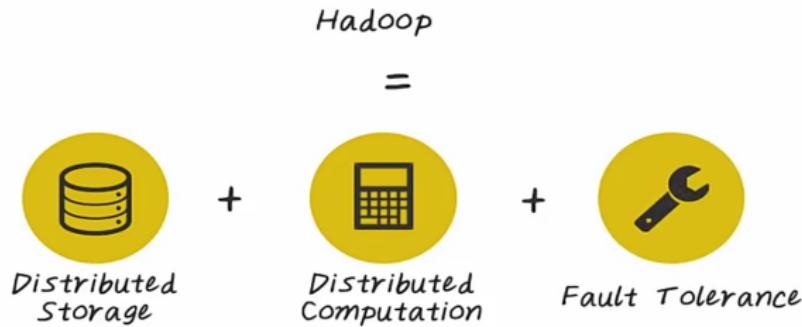
**Naive solution:** Run through the every data point one-by-one.

**Question:** Can we do faster?

**Another Question.** What happens if there are errors, failures during the computation process?

# Hadoop

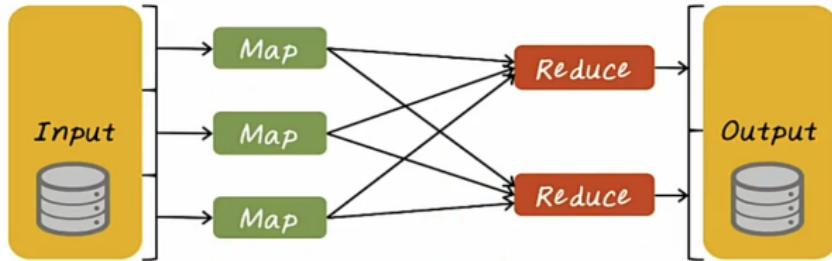
- ▶ Created in 2004 by Google, open source



Two key concepts in Hadoop

- ▶ **Map-Reduce** to process data
- ▶ **HDFS** to store data

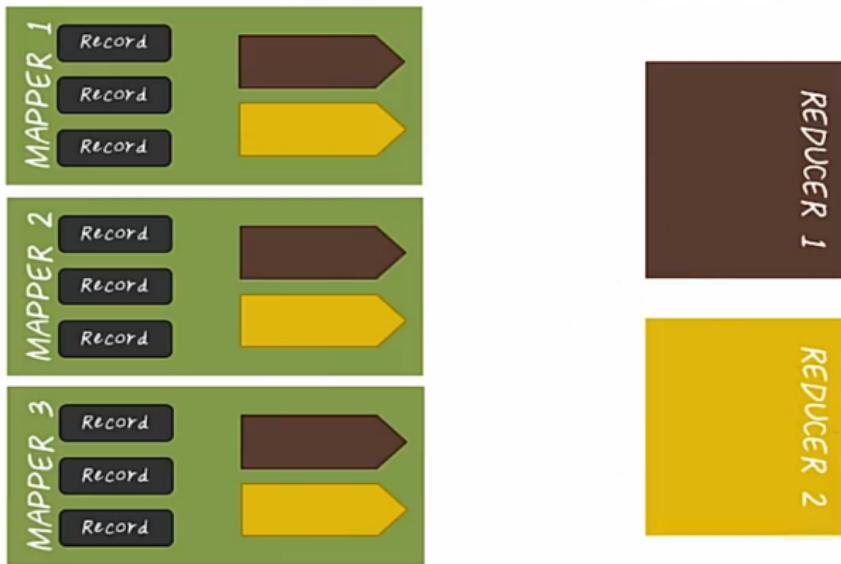
# Map-Reduce System



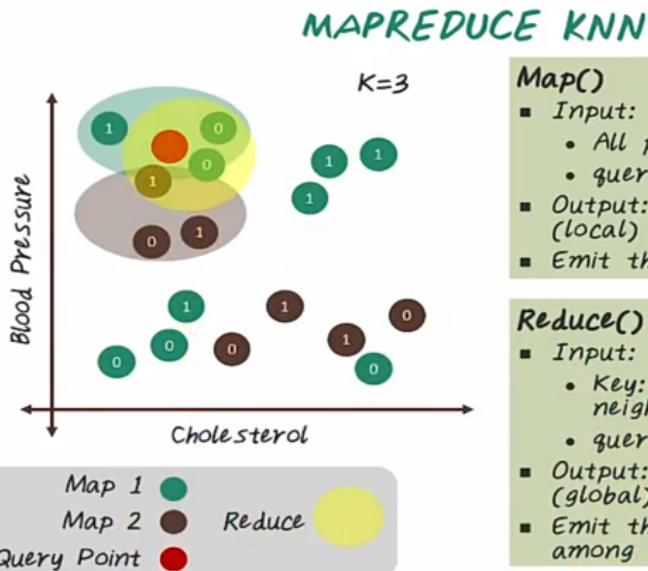
- ▶ Data flow in one direction (acyclic graph).
- ▶ Data flow from one disk to another disk inside the cluster.

# Map-Reduce System

## MAPREDUCE SYSTEM



# Example of Map-Reduce



## Map()

- Input:
  - All points
  - query point  $p$
- Output:  $K$  nearest neighbors (local)
- Emit the  $K$  closest points to  $p$

## Reduce()

- Input:
  - Key: null; values: local neighbors
  - query point  $p$
- Output:  $K$  nearest neighbors (global)
- Emit the  $K$  closest points to  $p$  among all local neighbors

# Hadoop distributed file system HDFS

(blackboard)

Main concepts:

- ▶ data node
- ▶ name node

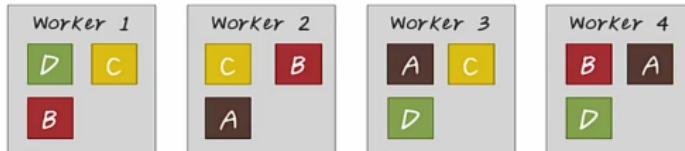
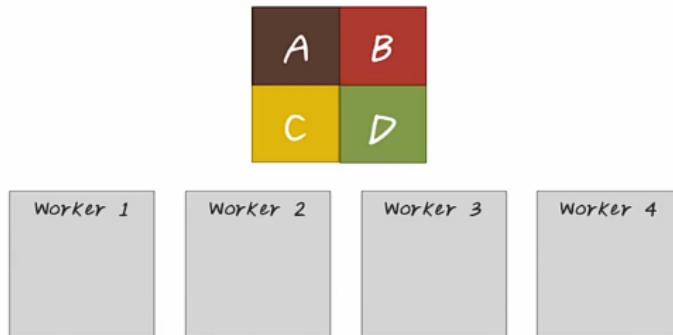
**Question:** Is there any problem in the case

- ▶ network failure
- ▶ disk failure on data node
- ▶ not all data node are used
- ▶ block sizes differ
- ▶ disk failure on name node

# Fault tolerance

**Question.** How to fix failure data node?

**Answer.** Duplicate every block to 3 nodes at random. Name node will control all the location of blocks.



## Fault tolerance

**Question.** What happens if name node dies?

**Answer.** All data are lost.

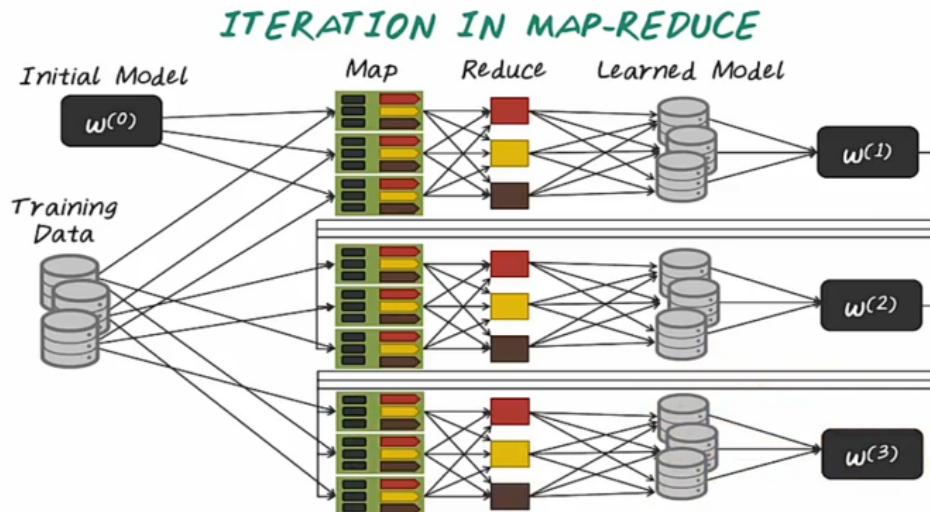
**Solution.** Create two name nodes: one *active* and one *standby*. If active dies, standby will become active and another name node is added.

## Benefits of Hadoop

1. Free (open-source)
2. Distributed Processing: data is stored in many different nodes
3. Fault Tolerance: default 3 replicas of each block are stored across the cluster.
4. Reliability: Automatically recover from failure.
5. High Availability: Data is available and accessible due to multiple copies of data.
6. Scalability: new hardware and new data can be easily added.
7. Data Locality: all computation is done and data only move locally in the cluster.

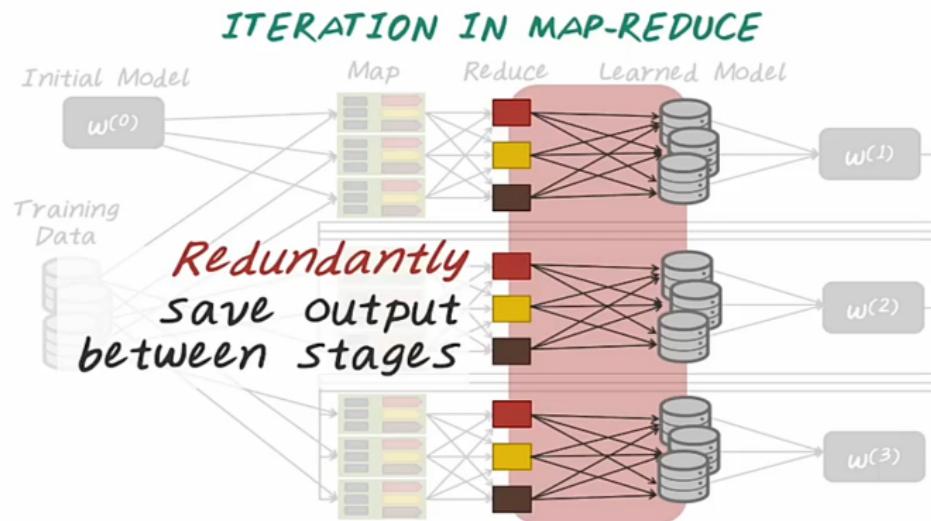
# Disadvantages of Hadoop

**Inefficient** for applications that repeatedly reuse data. Especially iterative algorithms (Machine learning, Graph Analysis).  
Example: *k*-means clustering algorithm.



# Disadvantages of Hadoop

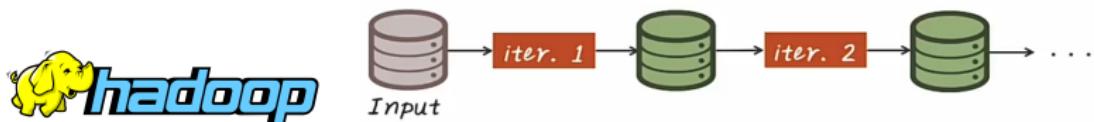
Example:  $k$ -means clustering algorithm.



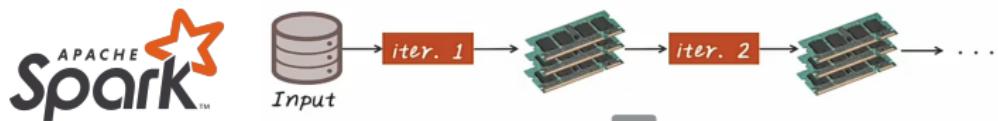
Recall previous week: The most time-consuming steps are reading/writing to hard disk, not computation.

# Solution: Spark

Instead of writing to hard drives at every step,

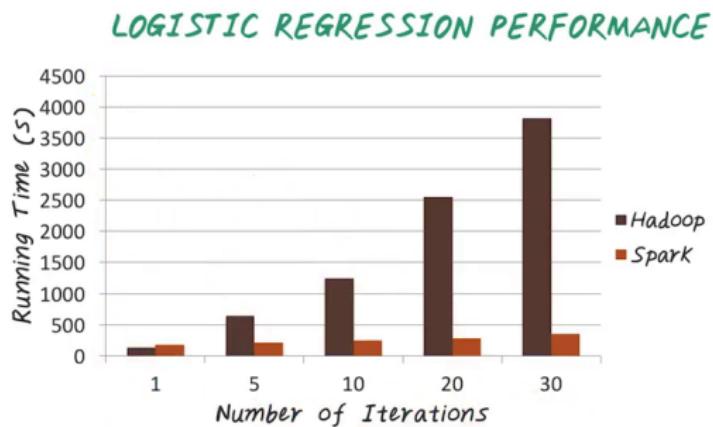


Spark keeps data in main memory



# Spark

- ▶ Created in 2012 by Berkeley University, open source
- ▶ Significantly faster than Hadoop



# Problems of Spark

**Problem 1.** RAMs are much more expensive than hard drives

**Problem 2.** RAMs are efficient but much *less stable* than hard drives →  
*less fault-tolerant*.

**Question.** How to make Spark fault-tollerant?

# Challenge

## CHALLENGE



Existing distributed storage abstractions depend on *fine-grained* updates

- Reads and writes to cells in a table
- E.g. databases, key-value stores, distributed memory

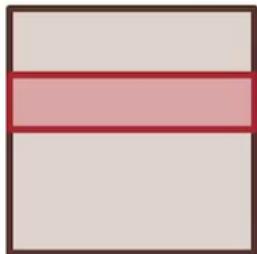
Require replicating data or logs across nodes for **fault tolerance** =



# Resilient Distributed Datasets (RDD)

RDD is the key concept in Spark.

## SOLUTION: RESILIENT DISTRIBUTED DATASETS (RDDS)



Provide an interface based on *coarse-grained* transformations  
(map, group-by, join, ...)

Efficient fault recovery using *lineage*

- Log one operation to apply to many elements
- Recompute lost partitions on failure
- No cost if nothing fails



# RDD Operations

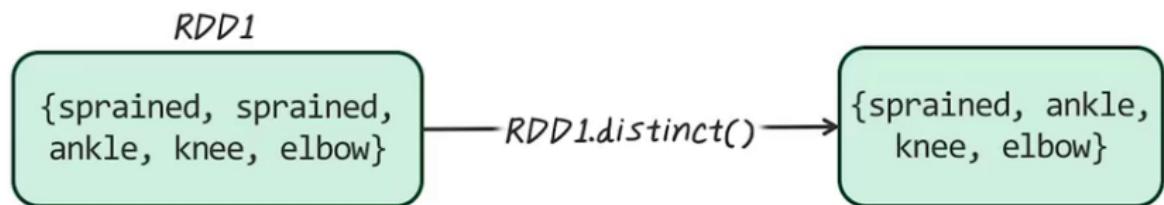
## SPARK OPERATIONS

<p><b>Transformations</b> (define a new RDD)</p>	<p>map filter sample groupByKey reduceByKey sortByKey</p>	<p>flatMap union join cogroup Cross mapValues</p>
<p><b>Actions</b> (return a result to driver program)</p>		<p>collect reduce Count save lookupKey</p>

# Some RDD Transformations

## RDD TRANSFORMATIONS

Operation: *Distinct()*

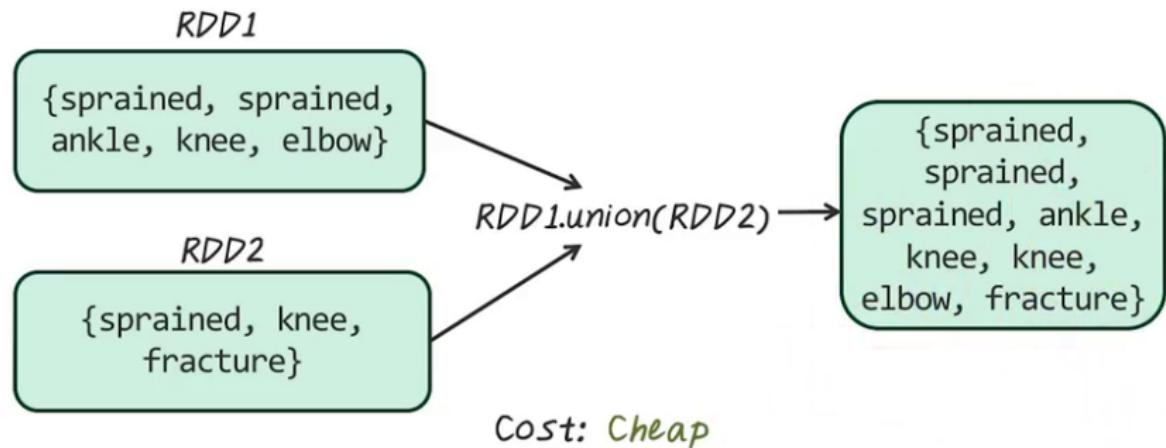


Cost: Cheap

# Some RDD Transformations

## RDD TRANSFORMATIONS

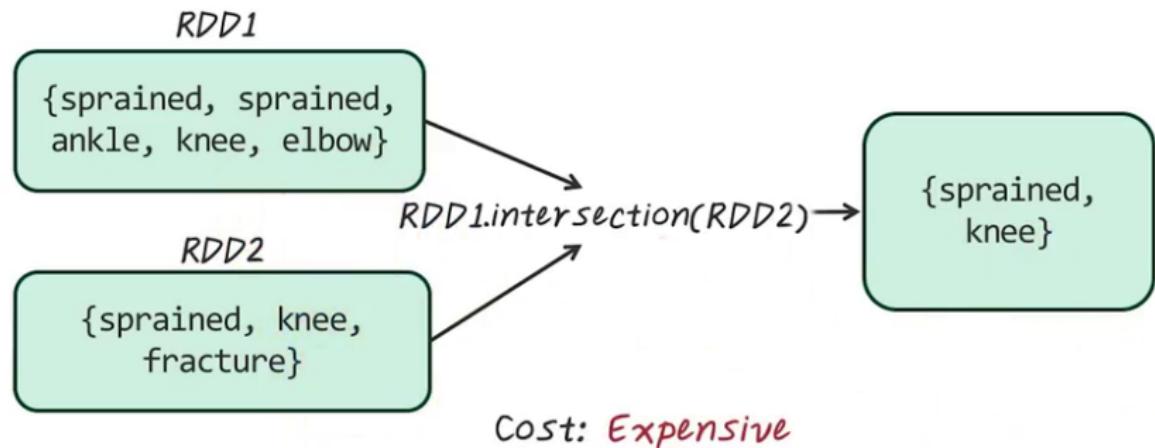
Operation: Union()



# Some RDD Transformations

## RDD TRANSFORMATIONS

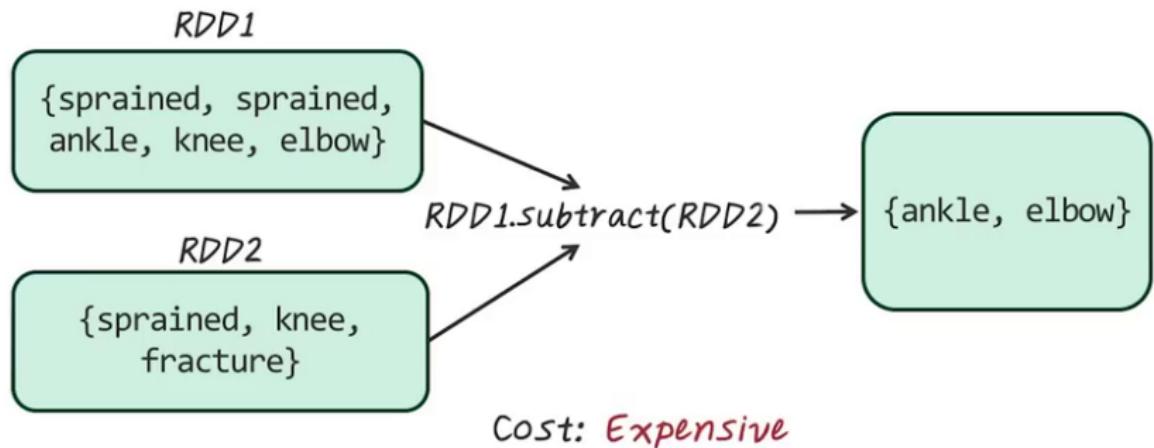
Operation: *Intersection()*



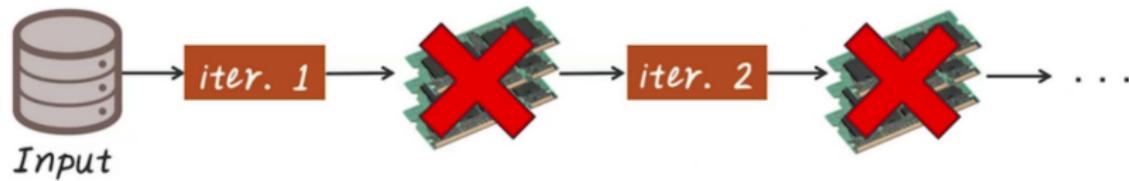
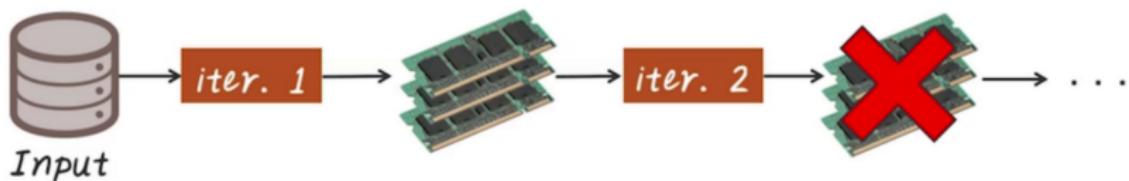
# Some RDD Transformations

## RDD TRANSFORMATIONS

Operation: Subtract()



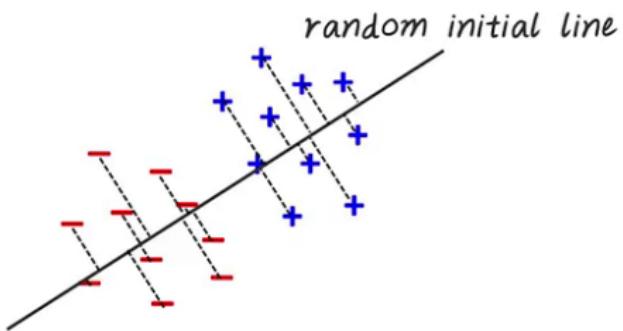
## RDD recovery



# Example of using Spark

## EXAMPLE: LOGISTIC REGRESSION

Goal: find best line separating two sets of points



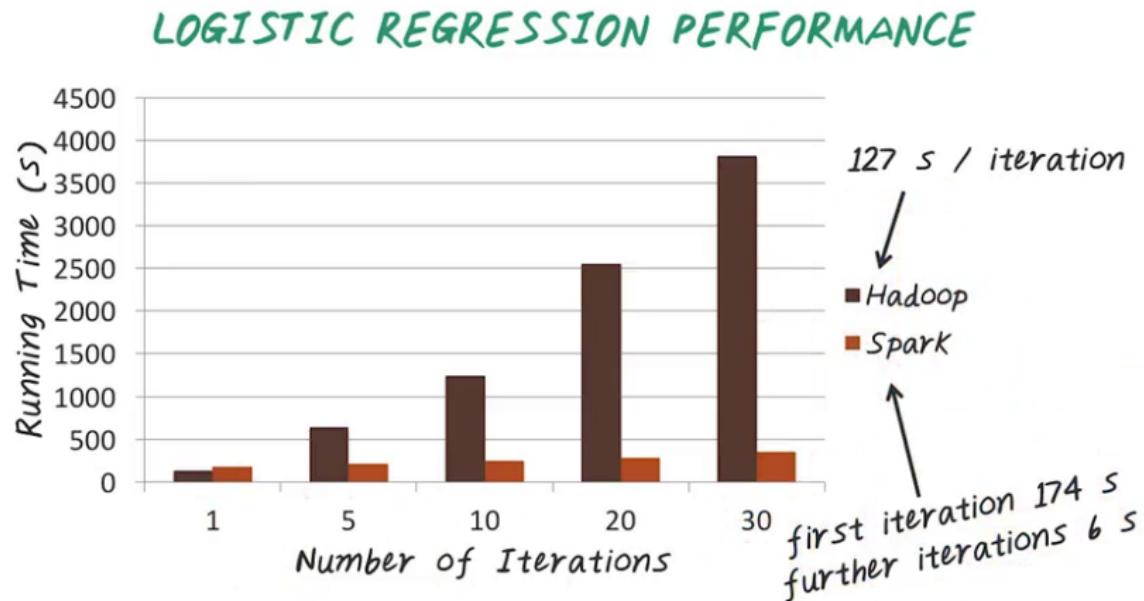
# Example of using Spark

## EXAMPLE: LOGISTIC REGRESSION

```
val data = spark.textFile(...).map(readPoint).cache()  
  
var w = Vector.random(D)  
  
for (i <- 1 to ITERATIONS) {  
    val gradient = data.map(p =>  
        (1 / (1 + exp(-p.y*(w dot p.x))) - 1) * p.y * p.x  
    ).reduce(_ + _)  
    w -= gradient  
}  
  
println("Final w: " + w)
```

Repeated MapReduce steps  
to do gradient descent

# Example of using Spark



# Conclusion

- ▶ Use Hadoop when
  - ▶ computation is one-way
  - ▶ budget is important
- ▶ Use Spark when
  - ▶ computation is iterative
  - ▶ speed is important

# Part II. Cloud Computing

# Contents

- 1. Cloud Computing**
- 2. Types of cloud services:** Public Cloud, Private Cloud, and Hybrid Cloud.
- 3. Levels of cloud services:** IaaS, PaaS, and SaaS.
- 4. Study cases**

# Why the cloud?

**In a nutshell:** With Netflix, who buys DVD anymore?

## Problems:

- ▶ Companies need to plan and build their new system in detail.
- ▶ Include: buildings and rooms, servers, networking devices, storage devices, extra security, cooling systems, power supplies.
- ▶ Usually overestimate or underestimate the computing needs for their new systems → waste of money.

## Ideal solution:

- ▶ Don't need to estimate cost and gambling financially on building the system
- ▶ Quickly grow and shrink your computing infrastructure based on your actual needs.
- ▶ Don't need to deal with security, cooling, software...

# Cloud Computing

**Cloud computing:** a pool of services such as storage, networking, and computers remotely via Internet.

## Benefit:

- ▶ Immediate access to various computing resources (networks, servers, storage, GPU, ...)
- ▶ Out of the box integration
- ▶ Elastic, on-demand, pay-per-use → reduce cost
- ▶ Minimal management effort: no security, no cooling...

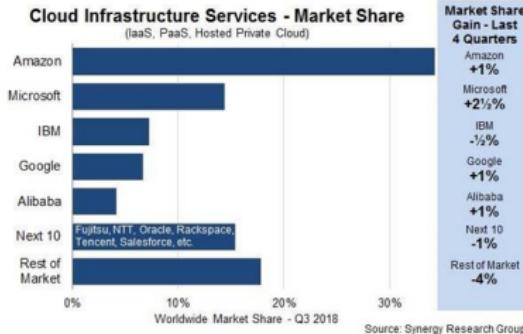
## Disadvantage:

- ▶ Need to trust cloud provider for your data.

# Cloud computing services

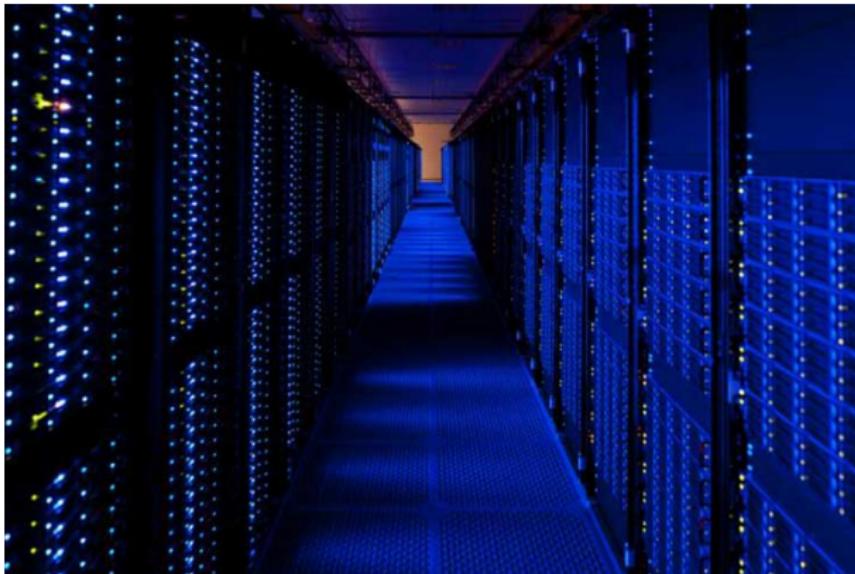


Google Cloud Platform



# Cloud Data Center

An Amazon cloud data center



# Technologies

Two key technologies: **virtualization** and high broadband **Internet**

Virtualization: a technique to share a single physical resources dynamically among multiple customers. Example:

- ▶ A big hard drive can be virtualized into several virtual small hard drives for different customers
- ▶ Run operation systems (Windows, Linux) on virtual machines (vmware, virtualbox) instead of installing directly on hardware.

VIRTUAL MACHINES



# Types of cloud services

1. Public cloud



2. Private cloud



3. Hybrid cloud



# Public Cloud

**Public cloud** is the most popular type of cloud model.

- ▶ Hardware is located at cloud provider, deliver via Internet
- ▶ Hardware is shared by several customers i.e. you and other customers may run on the same hardware.

## Benefit

- ▶ *Low costs*: pay only for the service they use.
- ▶ *No maintenance*: provider provides the maintenance.
- ▶ *Near-unlimited scalability*: On-demand resources.
- ▶ *High reliability*: A vast network of servers ensures against failure.

## Disadvantages

- ▶ *noisy neighbor effect*: performance fluctuations (CPU, RAM,...) depends on the use of other customers on the same machine.

# Private Cloud

**Private cloud:** To avoid noisy neighbor effect, customer want to get exclusive access to hardware.

**Location:** can be physically located at customer's or provider's site.

**Used by** government agencies, financial institutions, any other mid-to large-size business want to have better control.

## Benefits

- ▶ *Flexibility.* The organization can customize its cloud environment to meet specific business needs.
- ▶ *High security.* Resources are not shared with others, so higher levels of control and security are possible.
- ▶ *Good scalability.* Private clouds still afford the scalability and efficiency of a public cloud.

## Hybrid cloud

**Hybrid cloud:** data and applications can move between private and public clouds for greater flexibility and more deployment options.

### Example:

- ▶ Use the public cloud for high-volume, lower-security needs such as web-based email
- ▶ Use private cloud for sensitive, business-critical operations such as financial reporting.

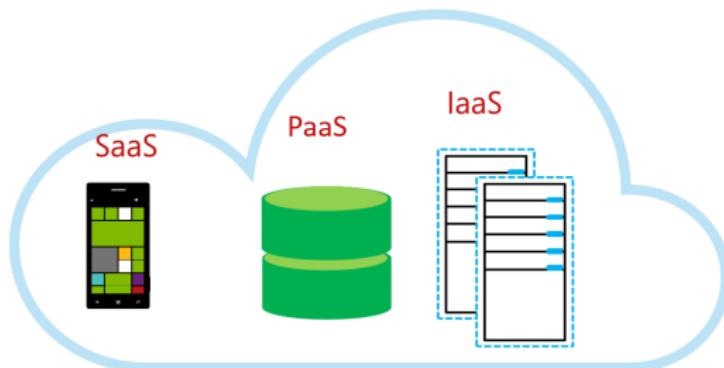
### Advantage:

- ▶ *Control.* Still control sensitive data.
- ▶ *Flexibility.* Access additional resources in the public cloud when needed.
- ▶ *Ease-of-use.* Transitioning between private and public cloud are easy.

# 3 levels of cloud services

## Cloud Service Types

- **SaaS – a complete software solution**
- **PaaS – a platform of services for hosting a custom solution**
- **IaaS – a way to run virtual servers in the cloud with full control**



# 3 levels of cloud services

## **Level 1. Software as a Service (SaaS)**

- ▶ Fully-formed software applications, delivered as cloud-based services.
- ▶ Customer subscribe to the service and use the application, normally through a web browser or by installing a client-side application.

**Example:** Office 365, Google Apps, Dropbox, Netflix

### **Advantage:**

- ▶ Easy access to applications without the need to install.
- ▶ No worry about issues such as updating applications and maintaining

# 3 levels of cloud services

## **Level 2: Platform as a Service (PaaS)**

- ▶ Provide resources and platform software on which developers can build their own applications.
- ▶ Include: operating system (OS), storage and compute capacity, and functional services for custom applications.

**Example:** Google App Engine, AWS Elastic Beanstalk, Windows Azure

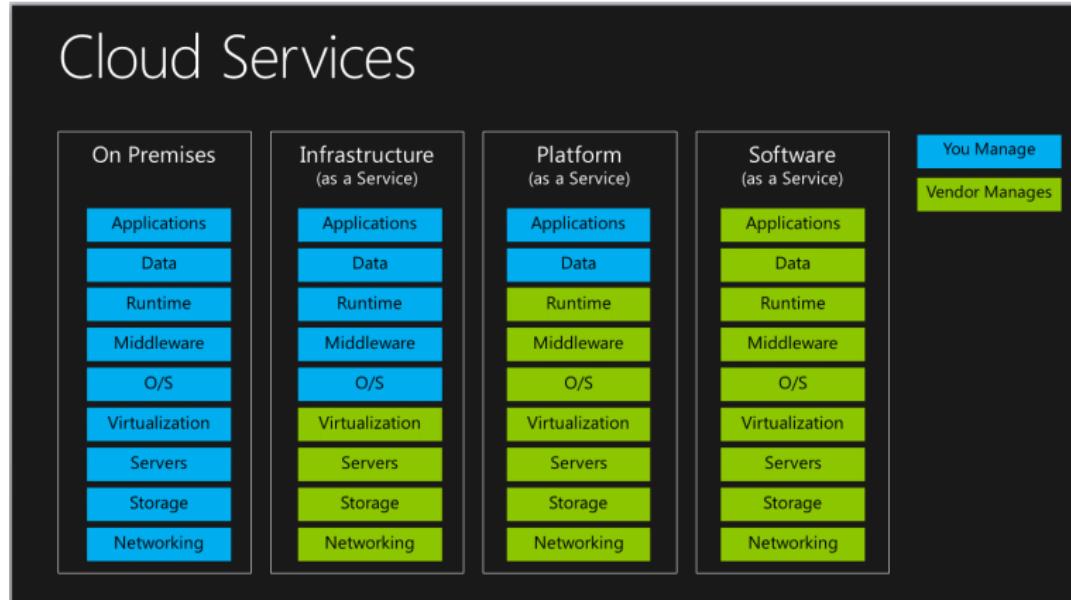
# 3 levels of cloud services

## **Level 3: Infrastructure as a Service (IaaS)**

- ▶ Provide virtualized server, network, and storage infrastructure components
- ▶ The maintenance of the OS and all software are up to customer.

**Example:** Amazon Web Services, Microsoft Azure

# 3 levels of cloud services



3 basic levels of cloud services: differ by which controlled by **cloud provider**, and which by the **customer**.

## Quick questions

**Q1.** You want to develop a PHP Web application using the PaaS cloud service. Who should be responsible for providing a web server with PHP interpreter and libraries installed?

- ▶ Provider
- ▶ Customer

**Q2.** You are using an IaaS cloud service with Windows virtual servers, and Microsoft has released a critical security patch. Who should install the patch on your server?

- ▶ Provider
- ▶ Customer

**Q3.** You are using a SaaS cloud service providing access to a Customer Relationship Management (CRM) package. Who is responsible for installing a new version of the package?

- ▶ Provider
- ▶ Customer

## Study case 1.

Suppose you want to have a set of

- ▶ 10 Linux systems with 4GB RAM each
- ▶ 2 Windows systems with 8GB each

to deploy your software.

Generally, a cloud provider

- ▶ creates the respective VMs in the background
- ▶ puts them in the same internal network
- ▶ provide you account, password → allowing you to access them

**Lead service:** Amazon EC2 package. Some benefits:

- ▶ secure and robust functionality.
- ▶ 99.99% uptime.
- ▶ specialized hardware for workloads: high graphics capability, high input/output (I/O), High Performance Computing (HPC).

## Study case 2.

Suppose you want to train your neural networks with deep learning

- ▶ Using high-end GPU
- ▶ only several hours per day.

**Lead service:** Amazon SageMaker package. Some benefits:

- ▶ up to 8 NVIDIA V100 GPU in parallel
- ▶ Most deep learning libraries available (TensorFlow, Keras, PyTorch...)
- ▶ Pay-per-minute use
- ▶ label data by human in demand (Amazon Mechanical Turk)

## Study case 3.

Suppose you want to have large storage ( $\geq 20$  TB) for your company (web server, big data, content distribution...)

**Lead service:** Amazon S3 package. Some benefits:

- ▶ fast access and high throughput
- ▶ boost access in Availability Zones: period with high frequent access
- ▶ several types: frequent access, infrequent access, archive (rarely access) with different prices

## Study Case 4.

Suppose you want to do computation on big data efficiently → using Hadoop or Spark.

**Lead service:** Amazon EMR package. Some benefits:

- ▶ Process data directly from Amazon S3
- ▶ Support Hadoop and Spark
- ▶ Amazon EMR Notebooks, based on Jupyter Notebooks, provide a managed environment for data scientists, analysts, and developers.

# Summary Cloud Computing

- 1. Cloud Computing**
- 2. Types of cloud services:** Public Cloud, Private Cloud, and Hybrid Cloud.
- 3. Levels of cloud services:** IaaS, PaaS, and SaaS.
- 4. Study cases**